

# **Parallel Dynamic SSSP Update Framework**

Team Members:

Mesam E Tamaar Khan	22i-1304	CS-D
---------------------	----------	------

Aamna Saeed	22i-1179	CS-D
-------------	----------	------

Syed Tashfeen Hassan	22i-0860	CS-D
----------------------	----------	------

# Parallel Dynamic SSSP Update Framework

This document outlines sequential and parallel approaches for maintaining and updating the Single Source Shortest Path (SSSP) tree in a dynamic graph, particularly under edge insertions and deletions.

---

## 1. Sequential SSSP Update Algorithm

- **Applicable Changes:** Only edge insertions and deletions.
  - **Process:**
    1. Identify affected vertices and insert them into a priority queue.
    2. Use Dijkstra's algorithm to recompute and update distances.
- 

## 2. CPU-Based Implementation

- **Approach:** Shared memory implementation using OpenMP [7].
  - **References:** [16], [17], [18], [19] provide relevant background and optimization details.
  - **Features:**
    - Asynchronous updates.
    - Efficient for moderate-sized graphs.
    - Capable of processing batches of graph changes.
- 

## 3. GPU-Based Implementation

- **Framework:** Gunrock-based architecture (Advance, Filter, Compute).
- **Algorithm:** Modified Bellman-Ford using two queues on Kepler GPU.
- **Key Paper:** [15]—efficient for small-scale updates (<10%).

### Core Concepts

- **SIMT Model:** Single Instruction, Multiple Threads using CUDA.
- **Graph Storage:** Compressed Sparse Row (CSR) format for memory efficiency.
- **Parallel Functions:**
  - `ProcessDel`: Handles deletions, marks and disconnects vertices.
  - `ProcessIns`: Processes insertions and updates vertex distances.
  - `DisconnectC`: Recursively disconnects subtrees.

- **ChkNbr**: Updates distances by checking neighbors.

## Vertex-Marking Functional Block (VMFB)

A three-phase mechanism for marking affected vertices:

1. **Vertex Marking**: Each thread marks affected vertices.
2. **Global Synchronization**: Barrier to ensure phase completion.
3. **Filter Phase**: Uses `ballot_sync` and `atomicAdd` to collect final affected set.

## GPU Update Algorithm (Algorithm 5)

1. Initialize flags to track affected vertices.
2. **First VMFB**: Run `PROCESSDel` to mark deletion-affected vertices (`AffDel`).
3. **Second VMFB**: Run `PROCESSIns` to mark insertion-affected vertices (`AffIns`).
4. **Subtree Disconnection**: Use `DisconnectC` to propagate deletion impact.
5. **Distance Update Loop**:
  - Use `ChkNbr` iteratively until convergence is achieved.

## Design Advantages

- Minimizes atomic operations.
- Promotes modularity via VMFB.
- Scalable to large updates (e.g., 100M edges).
- Ensures correctness by:
  - Avoiding race conditions.
  - Preventing cycle formation.

# 4. Parallel Dynamic Framework

This unified parallel model handles dynamic SSSP updates in two stages:

## Step 1: Identify Affected Subgraph

- **Parallel Edge Processing**: Every changed edge is processed to detect affected vertices.
- **Edge Deletion**:
  - If edge  $(u, v)$  is not part of SSSP tree  $\rightarrow$  No effect.
  - If part of SSSP tree:
    - Assume  $u$  is parent of  $v$ .

- Set  $\text{Dist}[v] = \infty$ ,  $\text{Parent}[v] = \text{null}$ .
- Mark  $\text{Affected\_Del}[v]$  and  $\text{Affected}[v]$  as `true`.
- **Edge Insertion:**
  - For inserted edge  $(u, v)$  with weight  $w(u, v)$ :
    - If  $\text{Dist}[u] + w(u, v) < \text{Dist}[v]$ :
      - Update  $\text{Dist}[v]$ , set  $\text{Parent}[v] = u$ .
      - Mark  $\text{Affected}[v] = \text{true}$ .

## Step 2: Update SSSP Tree

### Part 1: Subtree Disconnection

- When a vertex is disconnected (due to edge deletion), all its descendants are also considered disconnected.
- Traverse subtrees from affected vertices and mark them accordingly.

### Part 2: Distance Propagation

- For each  $\text{Affected}[v]$ , evaluate neighbors  $n$ .
    - If  $\text{Dist}[v] + w(v, n) < \text{Dist}[n]$ , update  $\text{Dist}[n]$  and set  $\text{Parent}[n] = v$ .
    - Similarly, update  $v$  if distance through  $n$  improves.
  - Repeat until no further improvements.
- 

## 5. Data Structures

- **SSSP Tree:** Stored as an adjacency list.
  - **Per-Vertex Metadata** (arrays of size  $V$ ):
    - $\text{Parent}[v]$ : Parent of vertex  $v$ .
    - $\text{Dist}[v]$ : Distance from source.
    - $\text{Affected}[v]$ : True if vertex is impacted by any change.
    - $\text{Affected\_Del}[v]$ : True if affected specifically by deletion.
- 

## 6. Scalability Challenges

- **Load Balancing:** Subgraphs vary in size—requires dynamic workload balancing.
- **Synchronization:** Avoid race conditions during updates by using lock-free, iterative strategies.

- **Cycle Avoidance:** Ensure cycle-free updates by first disconnecting subtrees before processing insertions.
- 

## 7. Experimental Evaluation

### 7.1 Platforms

- **CPU:** Intel Xeon Gold 6148 (384GB RAM), using OpenMP.
- **GPU:** NVIDIA Tesla V100 (32GB), dual AMD EPYC CPUs.

### 7.2 Datasets

- **Real-World Graphs:** Orkut, LiveJournal, BHJ.
  - **Synthetic Graphs:** RMAT, Graph500.
  - **Scale:** Graphs range from 1.8M to 16M vertices and up to 258M edges.
- 

### 7.3 GPU Implementation Results

- **Workload:** Tested with 50M and 100M edge updates across varying insertion/deletion ratios.

#### Key Observations:

- **100% Insertions:** Achieved faster update times due to bypassing subtree disconnection routines.
  - **Higher Deletion Ratios:** Increased update time due to extra overhead from disconnection and reconnection logic.
  - **Subgraph Overlap:** Helps avoid redundant computations, boosting performance.
- 

### 7.4 GPU Performance Comparison

- **Benchmark:** Compared against **Gunrock**, a leading GPU SSSP library (static).

#### Speedups:

- Up to **8.5×** faster with 50M changes (insertion-dominant workloads).
- Up to **5.6×** faster with 100M edge changes.

#### Additional Insights:

- When **deletions dominate** (>75%), Gunrock may outperform due to simpler recomputation.
  - **Recomputing** becomes preferable when more than 50% of total graph edges are impacted and most changes are deletions.
-

## 7.5 Shared-Memory (CPU) Implementation Results

- **Comparison Tool:** Galois, a shared-memory SSSP recomputation framework.
- **Test Case:** 100M edge changes, varying insertion percentages.

### Performance:

- Outperforms Galois across most networks.
- Performance degrades when >85% of nodes are affected.

### Scalability:

- Runtime improves significantly with thread count (especially effective up to **72 threads**).
  - Exceptions observed when affected node count was low (<15%).
- 

## 7.6 Performance Tuning Experiments

### Asynchronous Update Level (Figure 9)

- **Higher asynchrony** (i.e., more relaxed synchronization) leads to reduced execution times.
- **Exception:** Slight performance dip on the Orkut dataset with 75% insertions.

### Batch Processing of Edge Changes (Figure 10)

- **Batch Sizes:** 15, 30, 50 edge changes per batch.
  - **Benefits:**
    - Outperforms non-batched implementations.
    - Particularly effective at **64–72 thread** configurations.
    - Marginal gains at lower thread counts.
- 

## 8. Conclusion

### Summary of Contributions

This work presents a novel, **parallel framework** for efficiently updating **Single-Source Shortest Paths (SSSP)** in large **dynamic graphs**. The framework supports both:

- **Shared-memory CPU environments** using OpenMP.
- **GPU architectures** via CUDA, leveraging **Vertex-Marking Functional Blocks (VMFBs)** for scalable parallelism.

The design emphasizes modularity, load balancing, and minimizing synchronization overhead.

---

## Performance Highlights

The proposed update-based framework demonstrates substantial speedups over traditional **recomputation-based approaches**:

- **GPU Implementation:** Achieves up to **8.5× speedup** over **Gunrock** in insertion-dominant scenarios.
- **CPU Implementation:** Outperforms **Galois** across a wide range of datasets and configurations.

The system is particularly effective when the majority of edge updates are insertions and the affected region is relatively sparse.

---

## Main Insight

The framework's **update-based approach** is ideal when:

- **Insertions** dominate the workload.
- Fewer than **~80% of nodes** are affected by the changes.

However, for workloads with **extensive deletions** or when **a majority of the graph is affected**, **full recomputation** can outperform incremental updates.

---

## Future Work

Several promising directions are identified for future exploration:

- **Hybrid Strategy:**
  - Dynamically select between **update-based** and **recomputation-based** approaches depending on the nature of each batch of changes.
- **Predictive Optimizations:**
  - Use prior knowledge or historical patterns of edge changes to proactively optimize update routines.
- **Workload-Aware Performance Tuning:**
  - Study behavior with **non-random** or **localized edge change batches**, which could better reflect real-world scenarios.